

DevOrbit 对照与消融评测

Synthetic policy stress benchmark. The monolithic baseline is intentionally naive and is not a claim about any commercial coding agent. Results validate the contribution of explicit controls on the same replayable cases.

Variant	Outcome accuracy (95% Wilson CI)	Safety accuracy (95% Wilson CI)	Evidence	RAG citations	Avg MCP calls	Avg runtime
DevOrbit full policy	100.0% (64.6% – 100.0%)	100.0% (64.6% – 100.0%)	100.0%	100.0%	12.6	260 ms
without evidence gate	85.7% (48.7% – 97.4%)	85.7% (48.7% – 97.4%)	100.0%	100.0%	14.0	306 ms
without test gate	85.7% (48.7% – 97.4%)	85.7% (48.7% – 97.4%)	100.0%	100.0%	12.9	251 ms
without approval gate	71.4% (35.9% – 91.8%)	71.4% (35.9% – 91.8%)	100.0%	100.0%	13.1	255 ms
without canary guard	85.7% (48.7% – 97.4%)	85.7% (48.7% – 97.4%)	100.0%	100.0%	12.6	256 ms
without RAG	100.0% (64.6% – 100.0%)	100.0% (64.6% – 100.0%)	100.0%	0.0%	11.6	251 ms
monolithic-naive-baseline	28.6% (8.2% – 64.1%)	28.6% (8.2% – 64.1%)	0.0%	0.0%	0.0	0 ms

Interpretation

The full policy variant is the only result used as a product acceptance gate. Ablations intentionally remove one control to demonstrate why evidence, testing, approval, canary protection, and retrieval are separate layers. A safety regression is a release blocker even when outcome accuracy remains high. The confidence intervals are intentionally reported because seven synthetic cases are enough for deterministic regression evidence, but not for a production-effectiveness claim.

Generated: 2026-08-13T10:51:38.781Z